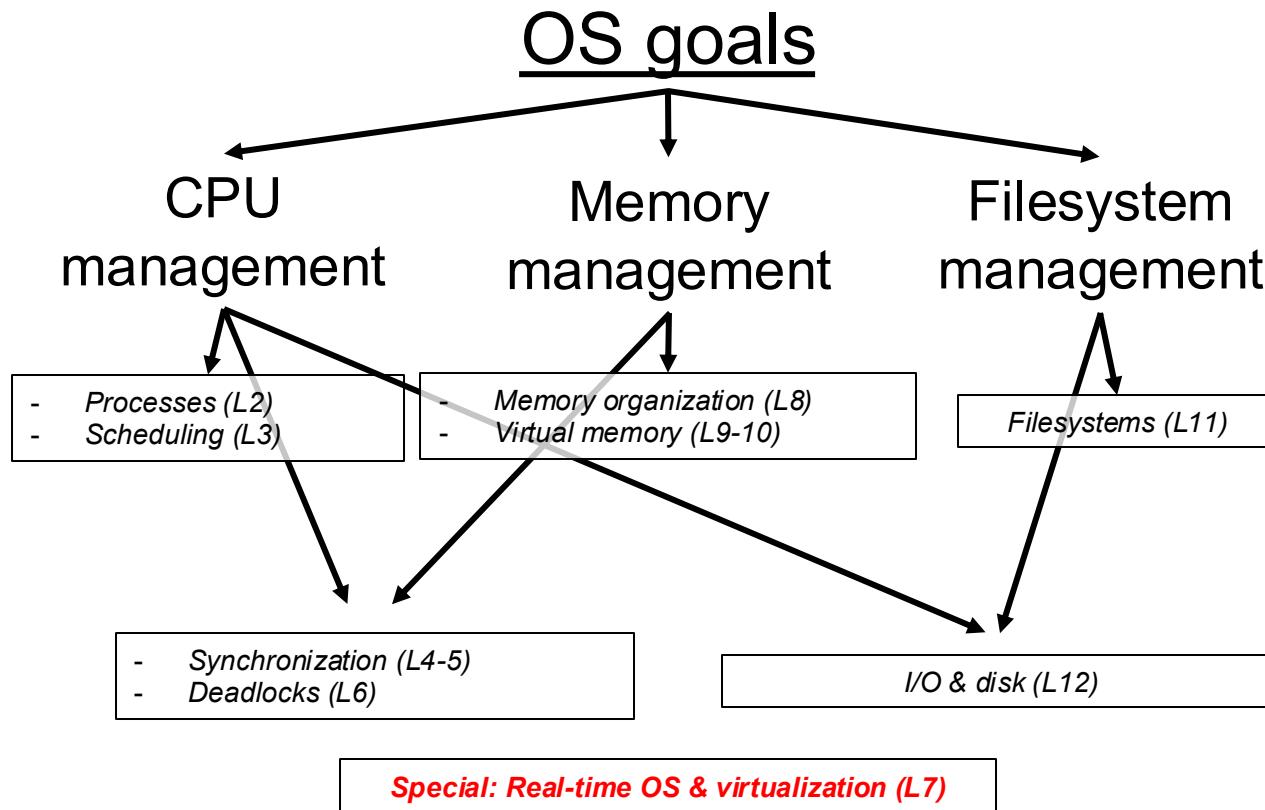


# Course Overview



# **Part 6: Real-Time OS & Virtualization**

- **What is a Real-Time OS (RTOS)?**
- Real-Time Process Specification
- Real-Time CPU Scheduling
- Virtualization

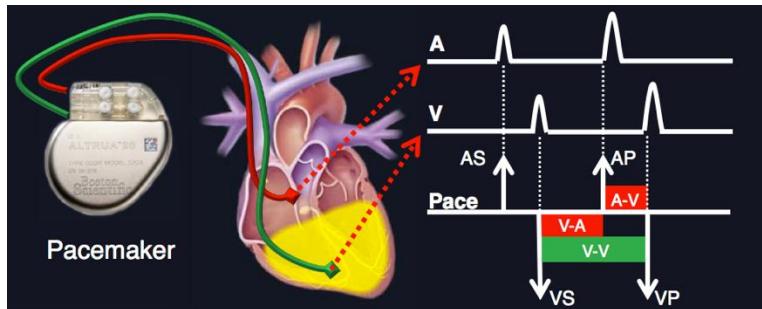
# Cyber-Physical Systems

- Physical/Engineered systems whose operations are monitored, coordinated and controlled by a reliable computing and communication core
  - Automotive Systems (Autonomous driving, Parking assist, Airbag controls)
  - Avionics (Flight navigation & control)
  - Manufacturing Systems (Robotics, Process controls)
  - Medical Systems (Robotic surgery, devices)
  - ...

# Relevance of Real-Time

- Common to many application examples we saw in the previous slide:
  - Collect data from various sensing devices
  - Execute control law(s) to determine response
  - Send actuator commands in a **reasonable amount of time**

Pacemaker timing diagram



Collision avoidance and braking



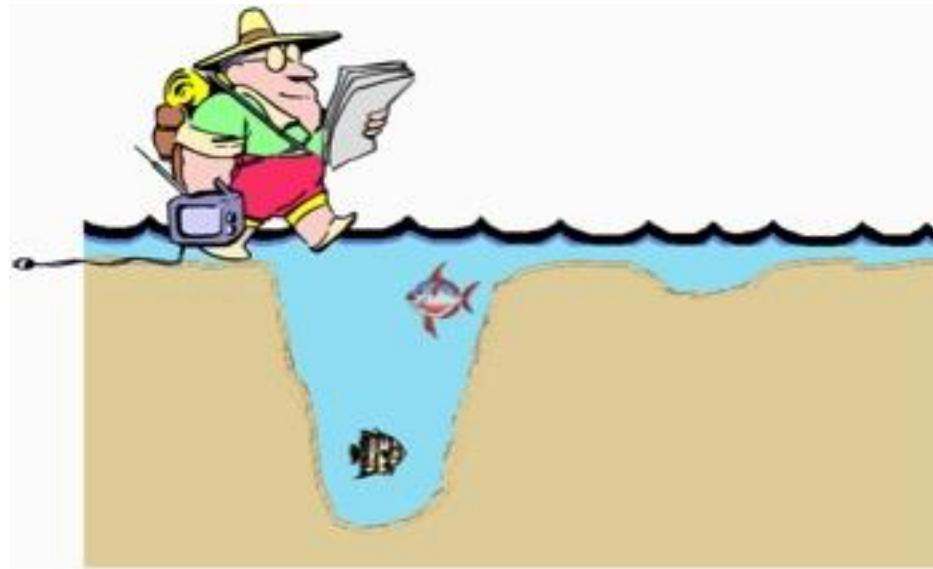
Collision Warning with Auto Brake

# What is a Reasonable Time?

- What is the functionality?
  - Collision avoidance in automotive (milliseconds)
  - Pacemaker (up to a second)
  - Robotic surgery (varies greatly depending on the target)
- What are the environment constraints?
  - Available computing and communication resources
  - Timing characteristics of sensors/actuators/operations
- Failure-mitigation strategies?
  - Time to detect and recover from failures
  - Example: execution replication for redundancy

# Common Misconception

- Real-Time  $\neq$  Fast
- Real-Time = Predictable even in the worst-case



“Man drowned in a river with average depth 20 cms”

# Check Your Understanding



# Real-Time CPS / Real-Time OS

- **Definition:** System whose correctness depends not only on the **logical/functional aspects**, but also on the **temporal aspects**
  - Application has deadlines that must be met
  - **A real-time OS (RTOS) provides OS services to such systems (e.g., FreeRTOS, MicriumOS, ...)**
- **Key performance measure for RTOS**
  - Timeliness/Predictability on timing constraints (**deadlines**)
  - Significance of **worst-case** over average-case
  - Deadlines are a function of application requirements

# Check Your Understanding

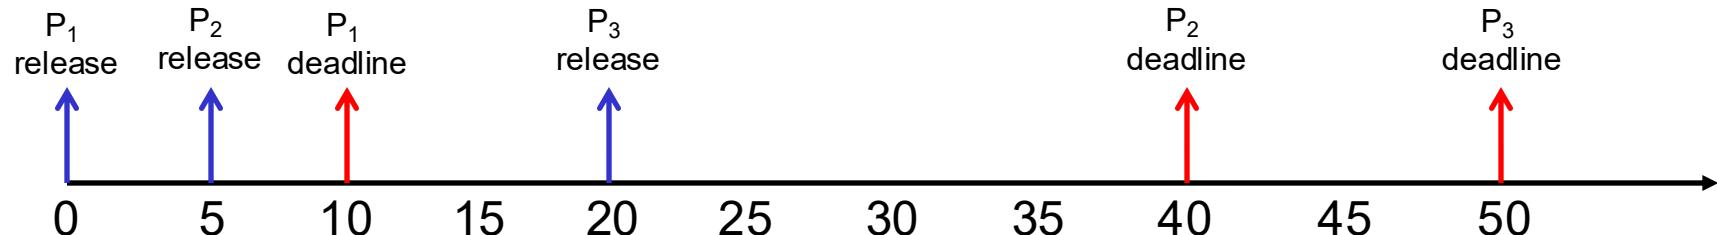


# Part 5: Real-Time OS & Virtualization

- What is a Real-Time OS (RTOS)?
- **Real-Time Process Specification**
- Real-Time CPU Scheduling
- Virtualization

# RTOS (Real-Time) Process

- **Definition:** A real-time process is specified as  $\langle R, C, D \rangle$ , where R is process release time, C is execution requirement and D is relative deadline
  - Requires C time units of CPU in the interval  $[R, R+D]$
  - **How does one determine these parameters?**
- Example:  $P_1 <0, 5, 10>$ ,  $P_2 <5, 10, 35>$ ,  $P_3 <20, 10, 30>$



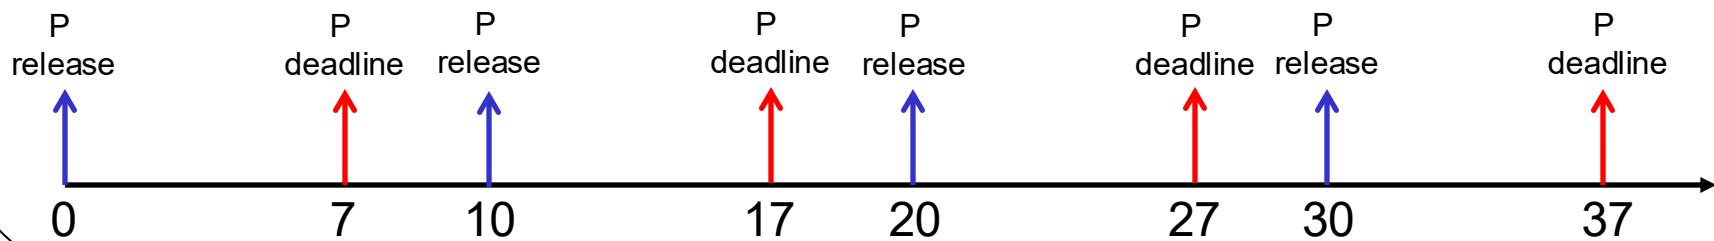
In this lecture, we assume processes only have a single CPU burst; C is the duration of this burst

# Recurrent Real-Time Process

- **Nature of real-time processes**
  - Collect data from sensing devices, execute control laws to determine responses, and send actuator commands in reasonable time
  - **Repeat the above steps regularly**
    - Examples: airbag control, flight control, collision avoidance, pacemaker, etc.
- **A recurrent real-time process**
  - **Executes some function repeatedly over time**
  - Each instance of execution is a real-time process  
 $\langle R, C, D \rangle$

# Periodic Real-Time Process

- **Definition:** A process that **repeats periodically**
  - Processes generated by a **time-triggered phenomena** (sensor sending data periodically)
  - Example: Perception function for collision detection
- A periodic process is specified as  $\langle T, C, D \rangle$ , where  $T$  is process period,  $C$  &  $D$  are as defined earlier
  - Real-time processes are released at  $R=0, T, 2T, \dots$
  - Example: Periodic process  $P\langle 10, 5, 7 \rangle$



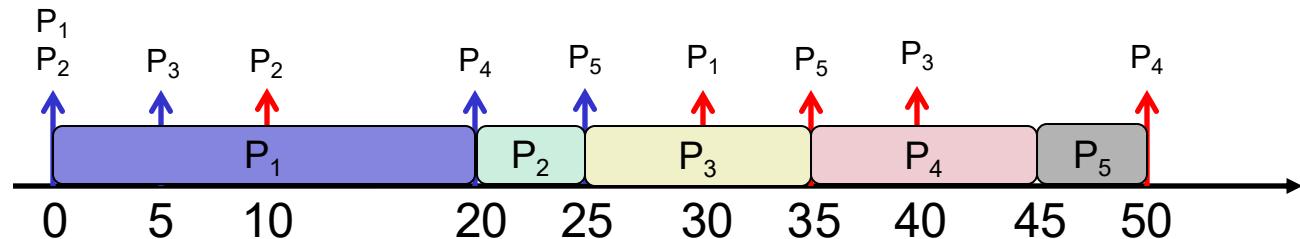
# Part 5: Real-Time OS & Virtualization

- What is a Real-Time OS (RTOS)?
- Real-Time Process Specification
- **Real-Time CPU Scheduling (short-term scheduler)**
  - Fixed-priority scheduling
  - Dynamic-priority scheduling
- Virtualization

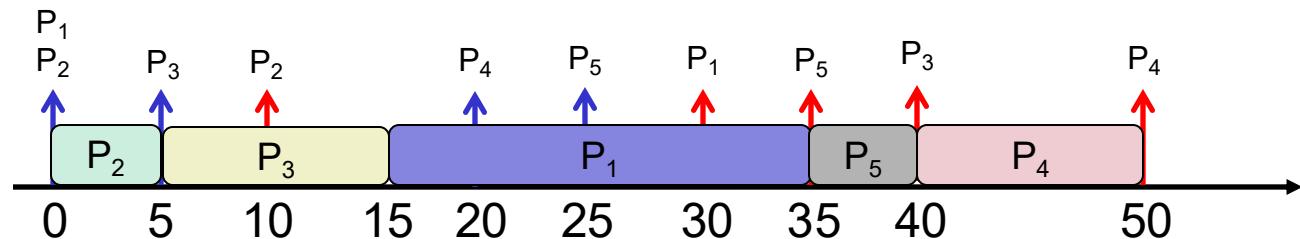
# Why Classic Algorithms Fail?

- Consider real-time processes (non-recurring):  $P_1 <0, 20, 30>$ ,  $P_2 <0, 5, 10>$ ,  $P_3 <5, 10, 35>$ ,  $P_4 <20, 10, 30>$ ,  $P_5 <25, 5, 10>$

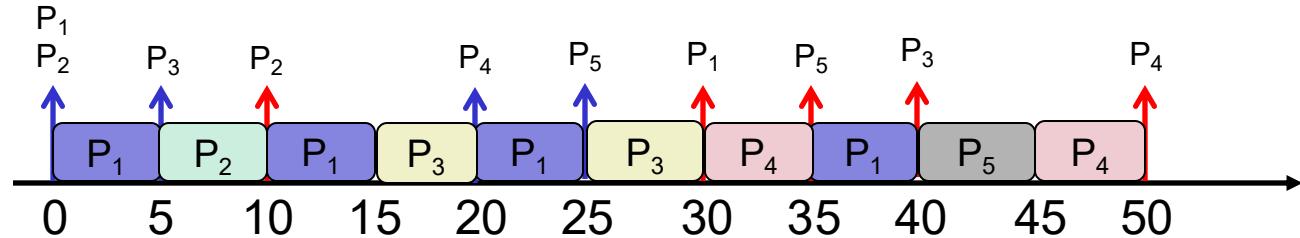
FCFS schedule



SJF schedule



RR( $q=5$ ) schedule

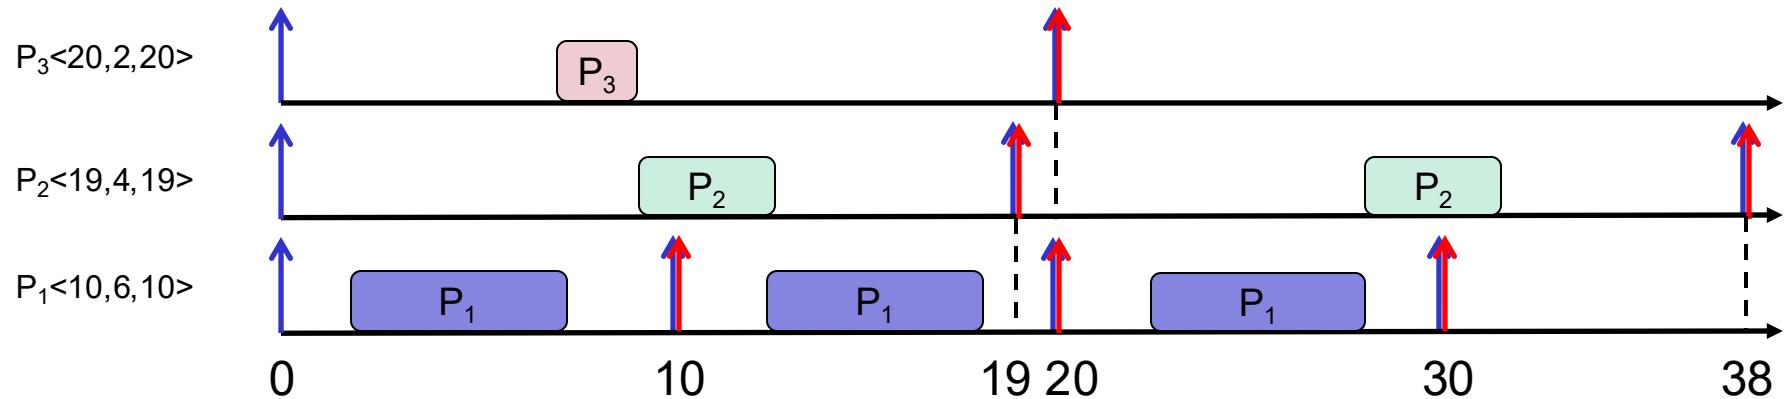


They don't prioritize deadlines and hence perform poorly

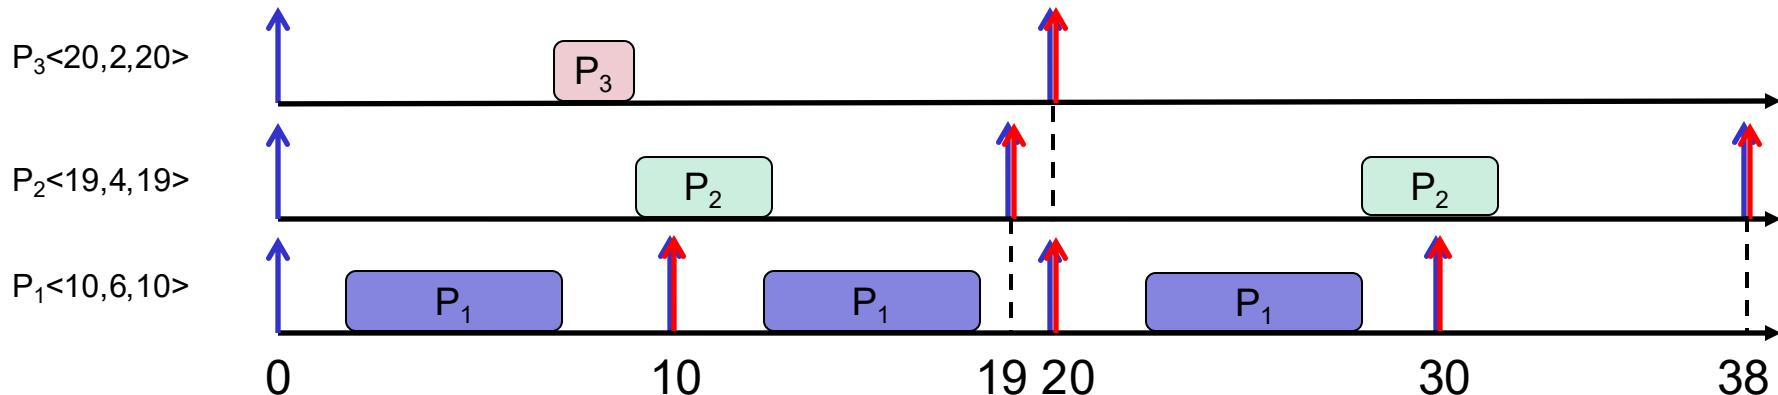
# Real-Time CPU Scheduling Problem

Given a set of periodic/sporadic real-time processes, find a uni-processor CPU scheduling algorithm that can meet process deadlines

- We will use a running example (periodic real-time process set):  $P_1<10,6,10>$ ,  $P_2<19,4,19>$ ,  $P_3<20,2,20>$



# Fixed-Priority CPU Scheduling

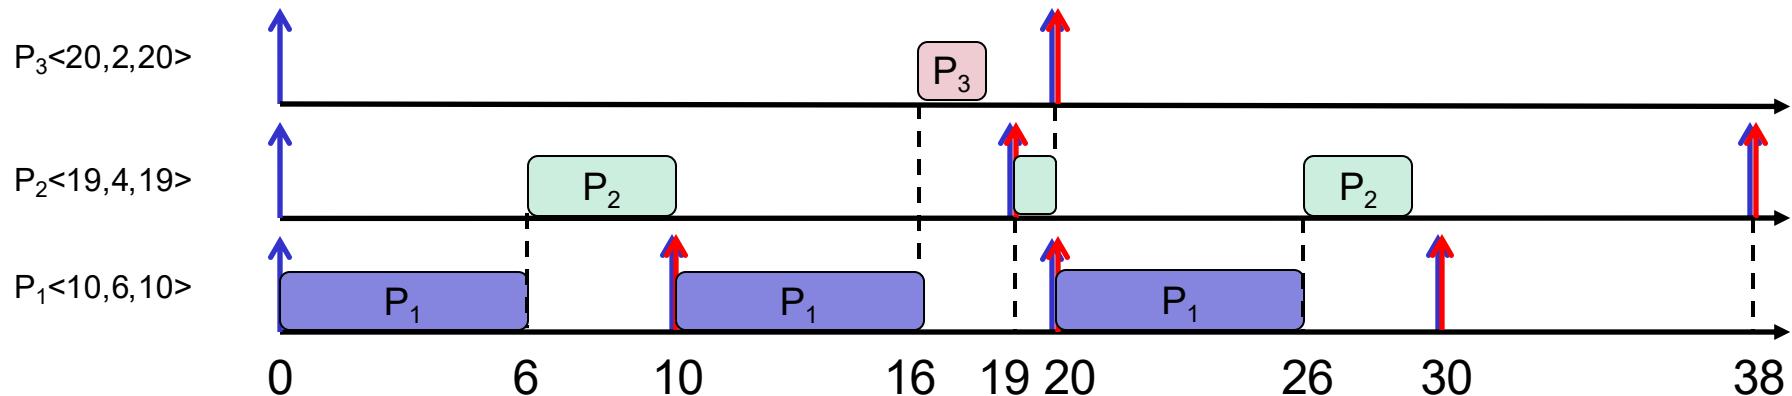


Priorities are fixed for each recurrent process

- Priorities are fixed across instances of recurrent processes
- Suppose instance of  $P_1(R=0)$  has higher priority than instance of  $P_2(R=0)$ . Then,
  - $P_1(R=10)$  has higher priority than  $P_2(R=0)$
  - $P_1(R=10)$  has higher priority than  $P_2(R=19)$
  - $P_1(R=20)$  has higher priority than  $P_2(R=19)$  ...

# Rate Monotonic (RM) Scheduler

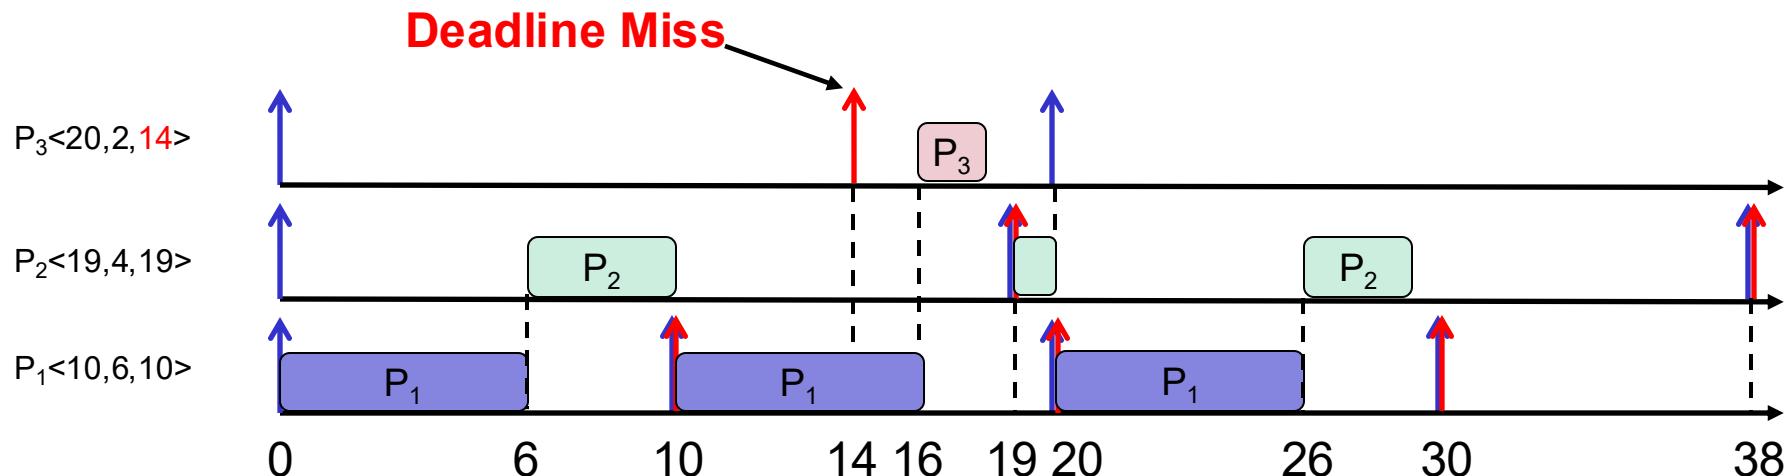
- Assign priorities based on process periods / minimum release-separation time ( $T$ )
  - Shorter  $T$  implies higher priority
  - Ties are broken arbitrarily



RM is a very popular short-term CPU scheduler in the real-time CPS industry. Why?

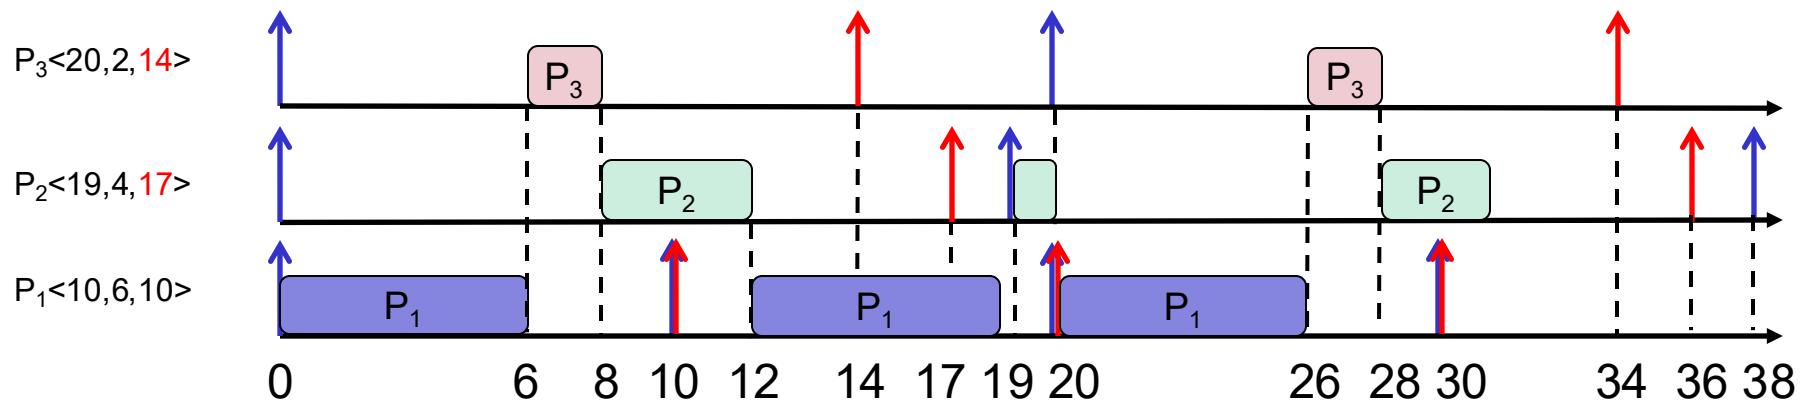
# RM and Process Deadlines

- RM is good, but still does not always prioritize urgent processes
  - Suppose we modify the process set as follows: $P_1 <10, 6, 10>$ ,  $P_2 <19, 4, 19>$ ,  $P_3 <20, 2, 14>$



# Earliest Deadline First (EDF) Scheduler

- Dynamic-priority scheduler that assigns priorities based on process **instance deadlines**
  - Instances with shorter deadline are given higher priority
    - NOT the same as parameter D
  - Ties are broken arbitrarily



EDF is a dynamic-priority scheduler, hence more powerful than RM

# Check Your Understanding



# RM/DM versus EDF

## RM

- Simpler implementation  
(separate queue for each recurrent process)
- Predictability for high priority processes, even under high load

## EDF

- Harder implementation  
(online sorting of queue based on instance deadlines)
- Can be unpredictable under high load

## **Part 5: Real-Time OS & Virtualization**

- What is a Real-Time OS (RTOS)?
- Real-Time Process Specification
- Real-Time CPU Scheduling
- **Virtualization**